

Sentinel Support: Full-Stack Fintech Case Resolution with Multi-Agent Automation

Build a production-minded **full-stack** system for internal support agents to (1) ingest and explore transactions, (2) generate **AI insights & reports**, and (3) **auto-resolve cases** (freeze card, open dispute, contact customer) via a **multi-agent** pipeline. The app must be **local/offline runnable**, have **deterministic fallbacks**, and expose **traces, metrics, and audit** artifacts.

Stack (required):

- **Frontend:** React + TypeScript (Vite or Next.js CSR mode), Tailwind (or CSS Modules)
- **Backend:** Node.js (TypeScript) + Express (or Fastify)
- **DB: PostgreSQL** (via Prisma/TypeORM/SQL)
- **Cache/Queue:** Redis (rate-limit + background jobs)
- **Infra:** Docker Compose (pg + redis + api + web)
- **Optional LLM:** behind a flag with deterministic fallback (no network required to pass)

1) What you're building (one sentence)

A **case-resolution console** where a support agent can load customer activity, get **AI-generated insights**, and run a **multi-agent triage** that recommends and executes safe actions, with **explainable traces, policy guardrails, and observability**.

2) Core Capabilities (must ship)

A) Frontend (React + TS)

- **Routes**
 - /dashboard: KPIs (alerts in queue, disputes opened, avg triage latency), quick filters.
 - /alerts: paginated/virtualized queue with risk score + "Open Triage".
 - /customer/:id: transactions timeline, category spend, merchant mix, anomalies.
 - /evals: run & view eval results (pass/fail, confusion matrix, top failures).
- **Triage Drawer (the hero)**
 - Shows **risk score, top reasons, plan, tool calls** (ok/error/duration), **fallbacks, citations, recommended action**.
 - Buttons: **Freeze Card, Open Dispute, Contact Customer, Mark False Positive**.
 - **Streaming updates** (SSE or WebSocket); drawer is fully **keyboard accessible**.

- **A11y & Perf**

- Focus trap, ESC to close, return focus; ARIA for dialog; polite live region for streamed updates.
- Virtualized tables for $\geq 2k$ rows; memoized rows; no jank.

B) Backend (Node + Express/Fastify + TS)

- **APIs**

- POST /api/ingest/transactions (CSV or JSON) → upsert, dedupe by (customerId, txnId).
- GET /api/customer/:id/transactions?from=&to=&cursor=&limit= → **keyset pagination**.
- GET /api/insights/:customerId/summary → categories, merchants, anomalies, monthly trend.
- POST /api/triage → starts a triage run (runId), streams events via **SSE** GET /api/triage/:runId/stream.
- POST /api/action/freeze-card (API key + optional OTP) → PENDING_OTP|FROZEN.
- POST /api/action/open-dispute → { caseId, status:"OPEN" }.
- GET /api/kb/search?q= → { results:[{docId,title,anchor,extract}] }.
- GET /metrics, GET /health.

- **Cross-cutting**

- **Rate-limit** 5 r/s per client (token bucket in Redis) → 429 with Retry-After.
- **Idempotency-Key** on ingest & actions; return prior result on replay.
- **Observability**: Prometheus metrics + structured JSON logs.
- **Audit**: every action appends case events (who/what/when, redacted payload).

C) Multi-Agent Orchestration (server-side)

- **Orchestrator (Planner)**: builds a bounded plan and executes sub-agents with timeouts & retries.
 - Default plan: ["getProfile", "recentTx", "riskSignals", "kbLookup", "decide", "proposeAction"]
- **Sub-Agents** (tool-using)
 - **Insights Agent**: categories, merchant concentration, spend patterns (deterministic rules; optional LLM phrasing).
 - **Fraud Agent**: velocity, device change, MCC rarity, prior chargebacks → {score, reasons[], action?}.
 - **KB Agent**: retrieve cited answers from local JSON (title + anchor).
 - **Compliance Agent**: OTP/identity gate, policy deny (e.g., unfreeze without verification).

- **Redactor:** PAN-like 13–19 digits → ****REDACTED****; mask emails; scrub logs & traces.
- **Summarizer:** customer message & internal note (template fallback).
- **Guardrails**
 - Tool **timeouts** ≤ 1s; flow budget ≤ 5s.
 - **Retries:** max 2 (150ms, 400ms + jitter).
 - **Circuit breaker:** open 30s after 3 consecutive failures per tool.
 - **Schema validation:** Zod/JSON-Schema for tool I/O (reject/annotate trace on mismatch).
 - **Prompt-injection:** user text cannot trigger tools without policy check; sanitize inputs.

3) Data Model (PostgreSQL)

```
-- customers/cards/accountscustomers(id pk, name, email_masked, kyc_level,
created_at)
cards(id pk, customer_id fk, last4, network, status, created_at)
accounts(id pk, customer_id fk, balance_cents, currency)
-- transactionstransactions(
  id pk, customer_id fk, card_id fk,
  mcc, merchant, amount_cents, currency,
  ts timestamptz, device_id, country, city
);-- indexes: (customer_id, ts DESC), (merchant), (mcc), (customer_id,
merchant)-- alerts & casesalerts(id pk, customer_id fk, suspect_txn_id fk,
created_at, risk text, status text)
cases(id pk, customer_id fk, txn_id fk null, type text, status text,
reason_code text, created_at)
case_events(id pk, case_id fk, ts, actor, action, payload_json)
-- triage runs & tracestriage_runs(id pk, alert_id fk, started_at, ended_at,
risk text, reasons jsonb, fallback_used bool, latency_ms int)
agent_traces(run_id fk, seq int, step text, ok bool, duration_ms int,
detail_json jsonb, primary key (run_id, seq))
-- kb & policieskb_docs(id pk, title, anchor, content_text)
policies(id pk, code, title, content_text)
```

4) Fixtures (must include)

```
/fixtures
  customers.json   cards.json   accounts.json
  transactions.json (≥ 200k rows; provide a generator script to reach ≥ 1M
locally)
  alerts.json      kb_docs.json policies.json chargebacks.json devices.json
  evals/*.json     (acceptance/golden cases; see §7)
```

5) Hard Requirements (SLOs/Safety)

- **Local run:** docker compose up brings everything online; seed scripts included.
- **Performance:** /customer/:id/transactions?last=90d **p95 ≤ 100ms** on a dataset ≥ 1,000,000 rows.
- **Streaming:** triage events stream to the FE (SSE or WS); resilient to reconnects.
- **Security:**
 - PAN-like sequences **never** appear in UI/logs/traces (13–19 digits redacted).
 - Mutations require X-API-Key; RBAC: “agent” vs “lead” (lead can force-approve); audit all actions.
 - **CSP** suitable for sensitive data pages (no unsafe-inline).
- **Correctness:** idempotent actions; 429 honored; evals pass (see §7).

6) APIs (contracts)

- POST /api/ingest/transactions → { accepted: true, count, requestId }
- GET /api/customer/:id/transactions?from=&to=&cursor=&limit=
 - **Keyset** pagination: returns { items:[], nextCursor }
- GET /api/insights/:customerId/summary
- POST /api/triage → { runId, alertId }
- GET /api/triage/:runId/stream (SSE) → events: plan_built, tool_update, fallback_triggered, decision_finalized
- POST /api/action/freeze-card (Idempotency-Key, API key)
 - Req: { cardId, otp? } → { status: "PENDING_OTP" | "FROZEN", requestId }
- POST /api/action/open-dispute
 - Req: { txnId, reasonCode, confirm } → { caseId, status: "OPEN" }
- GET /api/kb/search?q=
- GET /metrics, GET /health

7) Acceptance Scenarios (we will verify exactly)

1. Freeze w/ OTP path

- a. From /alerts, open triage on an alert.
- b. **Expected:** recommendation = “Freeze Card”, **OTP required**; enter 123456 → status FROZEN; trace shows freezeCard ok; metrics increment action_blocked_total{policy=otp_required}.

2. Dispute creation

- a. Unrecognized ₹4,999 at "ABC Mart" yesterday.
 - b. **Expected:** single txn match; reason **10.4** proposed; **case OPEN** with caseId in timeline; citation to KB **Disputes**.
- 3. **Duplicate pending vs captured**
 - a. "Charged twice at QuickCab."
 - b. **Expected:** explanation (preauth vs capture); **no dispute**; risk = low|medium; KB citation appears.
- 4. **Risk tool timeout → fallback**
 - a. Simulate riskSignals failure.
 - b. **Expected:** fallback_used=true, risk ≤ medium, reason includes risk_unavailable; SSE shows fallback_triggered.
- 5. **429 behavior**
 - a. Trigger rate limit (spam triage).
 - b. **Expected:** 429 with Retry-After; FE disables control ~given ms; **no duplicate triage runs**; next attempt succeeds.
- 6. **PII redaction**
 - a. User input or tool detail contains 4111111111111111.
 - b. **Expected:** redacted as ****REDACTED**** in UI/logs/trace; structured log line includes masked=true.
- 7. **Performance**
 - a. With ≥1M txns, /customer/:id/transactions?last=90d **p95 ≤ 100ms**; show timing + EXPLAIN ANALYZE snippet in README.

8) Metrics & Logs (what we expect to see non-zero)

- **Metrics**
 - api_request_latency_ms (hist)
 - agent_latency_ms (hist)
 - tool_call_total{tool,ok}
 - agent_fallback_total{tool}
 - rate_limit_block_total
 - action_blocked_total{policy}
- **Structured logs (JSON)**
 - fields: ts, level, requestId, runId, sessionId, customerId_masked, event, masked=true
 - key events: plan_built, tool_invoked, fallback_triggered, decision_finalized, action_completed

9) Evals (golden set & CLI)

- Provide **≥12 eval cases** in `/fixtures/evals/*.json`—cover all acceptance scenarios + ambiguous merchants + device change + travel window cases.
- CLI `npm run eval` prints:
 - **Task success rate, fallback rate** by tool
 - **Agent latency p50/p95**
 - **Risk confusion matrix** (low|medium|high)
 - **Top policy denials**

10) Deliverables

- Monorepo (or two folders): `/web` (React), `/api` (Node), `/fixtures`, `/scripts`, `/docs`.
- **Docker Compose** for pg, redis, api, web.
- **README (≤1 page)**: run in ≤3 commands, arch diagram (ASCII ok), key trade-offs.
- **ADR.md**: 8–12 bullets (why keyset pagination, why SSE, why schema X, etc.).
- **Postman/HTTP** collection.
- **Demo video (≤8 min)**: show freeze (OTP), dispute, fallback, 429, metrics/logs.
- **Eval report** (from CLI output) + brief commentary.

```
{ "topMerchants": [{ "merchant": "ABC", "count": 12 } ], "categories": [{ "name": "Transport", "pct": 0.23 } ], "monthlyTrend": [{ "month": "2025-07", "sum": 120045 } ], "anomalies": [{ "ts": "2025-07-13", "z": 3.1, "note": "spike" } ] }
```